

Distributed and Parallel Computing

Lecture 02

Alan P. Sexton

University of Birmingham

Spring 2018

Measuring Parallel Speedup

- **Latency:** the time from initiating to completing a task
 - Units of time
- **Work:** a measure of what the CPU/GPU has to do for a particular task
 - number of floating point operations
 - number of images processed
 - number of pixels processed
 - number of simulation steps
- **Throughput:** work done per unit time

Speedup and Efficiency

- **Speedup_P, S_P** : The ratio of the latency for solving a problem with 1 hardware unit to the latency of solving it with P units
 - $S_P = \frac{T_1}{T_P}$
 - Perfect linear speedup: $S_P = P$
- **Efficiency, E_P** : The ratio of the latency for solving a problem with 1 hardware unit to P times the latency of solving it on P units
 - This measures how well the individual hardware units are contributing to the overall solution
 - $E_P = \frac{T_1}{P \times T_P} = \frac{S_P}{P}$
 - Perfect linear efficiency: $E_P = 1$

Interpreting speedup and efficiency

- Sub-linear speedup and efficiency is normal
 - Overhead in parallelizing a problem
- Super-linear speedup is possible, but usually due to special conditions
 - e.g. Serial version does not fit in CPU cache but each of the parallel sub-problems do.
- Important to compare the best serial version of the program with the parallel version
 - Serial algorithm A is fast but hard to parallelize
 - Serial algorithm B is slow but gives Parallel algorithm B
 - To measure speedup/efficiency, compare Serial A to Parallel B
 - Both must solve the same problem, but allow for minor differences
 - e.g. small round-off differences (but be aware of differing precision on host and on GPU!)
 - differences due to different execution order

Strong Scalability

Gene Amdahl, in 1967, argued that the time spent executing a program is composed of the time spent doing non-parallelizable work plus the time spent doing parallelizable work:

$$T_1 = T_{\text{ser}} + T_{\text{par}}$$

Therefore, if the speedup on P units of the parallel part only is s

$$T_P = T_{\text{ser}} + \frac{T_{\text{par}}}{s}$$

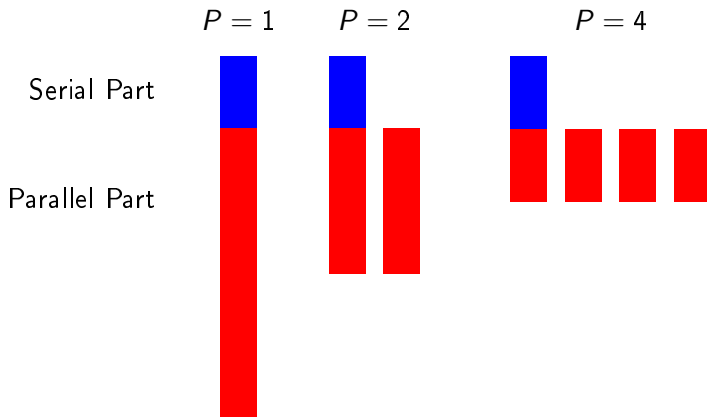
Hence the overall speed up given the speedup of the parallel part is s is:

$$S_P = \frac{T_{\text{ser}} + T_{\text{par}}}{T_{\text{ser}} + \frac{T_{\text{par}}}{s}}$$

If we let f be the fraction of a program that is parallelizable, then $T_{\text{ser}} = (1 - f)T_1$ and $T_{\text{par}} = fT_1$. Hence

$$S_P = \frac{(1 - f)T_1 + fT_1}{(1 - f)T_1 + \frac{fT_1}{s}} = \frac{1}{1 - f + \frac{f}{s}} \quad (\text{Amdahl's Law})$$

Amdahl's Law Graphically



Interpreting Amdahl's Law

Amdahl's law seems to say that there is a limit to parallel speedup

$$\lim_{s \rightarrow \infty} S_P = \lim_{s \rightarrow \infty} \frac{1}{1 - f + \frac{f}{s}} = \frac{1}{1 - f}$$

or, in other words

$$\begin{aligned} \lim_{s \rightarrow \infty} \frac{T_1}{T_P} &= \frac{1}{1 - f} \\ \Rightarrow \lim_{P \rightarrow \infty} T_P &= T_{\text{ser}} \quad \text{assuming } P \rightarrow \infty \Rightarrow s \rightarrow \infty \end{aligned}$$

John Gustafson and Edwin Barsis, in 1988, argued that Amdahl's law did not give the full picture

- Amdahl kept the task fixed and considered how much you could shorten the processing time by running in parallel
- Gustafson and Barsis kept the processing time fixed and considered how much larger a task you could handle in that time by running in parallel
- This was motivated by observing that, as computers increase in power, the problems that they are applied to often increase in size

Assume that W is the workload that can be executed without parallelism in time T . If f is the fraction of the workload that is parallelizable, then

$$W = (1 - f)W + fW$$

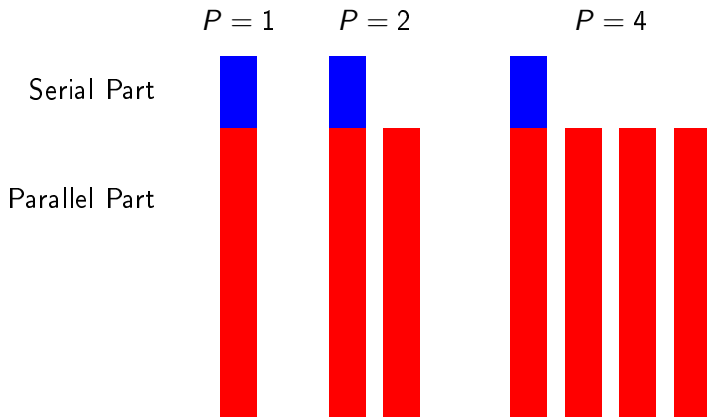
With speedup s , we can run s times the parallelizable part in the same time, although we don't change the amount of work done in the non-parallelizable part:

$$W_s = (1 - f)W + sfW$$

If we do W_s in time T , we are, on average, doing W amount of work in time $\frac{TW}{W_s}$. Hence the total speedup is:

$$S_s = \frac{T}{TW/W_s} = \frac{W_s}{W} = 1 - f + fs \quad (\text{Gustafson-Barsis})$$

Gustafson-Barsis Law Graphically



- Both Amdahl (A) and Gustafson-Barsis (GB) are correct
- The two together give guidance on what kinds of problems can benefit from parallelization and how
- You can only go so much faster on a fixed problem by parallelization
- You can avoid Amdahl's limit on speedup if you can increase the size of the parallelizable part of the problem faster than you increase the size of the non-parallelizable part