

# Machine Learning, Machine Learning (extended)

## 7 – Unsupervised Learning: Clustering

Kashif Rajpoot

[k.m.rajpoot@cs.bham.ac.uk](mailto:k.m.rajpoot@cs.bham.ac.uk)

School of Computer Science

University of Birmingham

# Outline

- Supervised vs unsupervised learning
- Clustering
- Similarity measure
- K-means clustering
- Kernelized k-means clustering
- Hierarchical agglomerative clustering

# Supervised vs unsupervised learning

- So far, we have looked at supervised learning
- Supervised learning
  - The algorithm learns from  $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N$  and  $t_1, t_2, \dots, t_N$  so that it can later classify  $\mathbf{x}_{new}$
  - The availability of  $t_n$  makes learning a supervised task
- What if we only have  $\mathbf{x}_n$ , but not  $t_n$ ?
  - Unsupervised learning
- Clustering: create a grouping of objects
  - Each group containing “similar” objects


# Clustering: examples

- Shopping recommendation system
  - Recommend products likely to be purchased by a customer
- Grouping people from social network
  - Based on user activity
- Brain region network analysis
  - Determine “similar” brain regions that “activate” together or “rest” together

# Clustering


- Clustering: partition data into groups such that each group contains “similar” objects
  - Increase intra-group similarity
  - Similar objects should belong to same cluster
- How do we measure “similarity” between objects?
  - “closeness”

$$\begin{aligned} D_{ij} &= (\mathbf{x}_i - \mathbf{x}_j)^T (\mathbf{x}_i - \mathbf{x}_j) \\ &= \left(x_i^{(1)} - x_j^{(1)}\right)^2 + \left(x_i^{(2)} - x_j^{(2)}\right)^2 \\ &= \sum_{d=1}^M \left(x_i^{(d)} - x_j^{(d)}\right)^2 \end{aligned}$$



# Clustering: Similarity measure


- Clustering: partition data into groups such that each group contains “similar” objects
  - Increase intra-group similarity
  - Similar objects should belong to same cluster
- How do we measure “similarity” between objects?
  - “closeness”
- Which points are “closest” or “most similar”?
- Choice of similarity measure is dependent on the nature of data and your application




# Clustering: Similarity measure

- How to estimate the distance (i.e. similarity) between two objects/points  $\mathbf{p}$  and  $\mathbf{q}$ ?
  - Various measures can be used
- K-means algorithm is flexible to use any distance measure


Euclidean distance



Manhattan distance



Cosine angle



$$D_e(\mathbf{p}, \mathbf{q}) = \sqrt{\sum_{d=1}^M (p^{(d)} - q^{(d)})^2}$$


$$D_m(\mathbf{p}, \mathbf{q}) = \sum_{d=1}^M |p^{(d)} - q^{(d)}|$$

$$D_c(\mathbf{p}, \mathbf{q}) = \cos^{-1} \left( \frac{\sum_{d=1}^M p^{(d)} q^{(d)}}{\sqrt{\sum_{d=1}^M p^{(d)2}} \sqrt{\sum_{d=1}^M q^{(d)2}}} \right) / \pi$$

# K-means clustering

- Consider that there are total  $K$  clusters
- In 2D, each cluster  $k$  is represented by a cluster centre (i.e. mean)

$$\boldsymbol{\mu}_k = [\mu_k^{(1)}, \mu_k^{(2)}]^T$$



- Each object  $\mathbf{x}_n$  is assigned to its closest cluster  $k$  on the basis of its distance to cluster mean  $\boldsymbol{\mu}_k$
- Distance to cluster  $k$  can be estimated in various ways
  - For example: squared Euclidean distance

$$D_{nk} = (\mathbf{x}_n - \boldsymbol{\mu}_k)^T (\mathbf{x}_n - \boldsymbol{\mu}_k)$$

- No analytical solution for finding the closest cluster mean  $\boldsymbol{\mu}_k$  for all objects
  - Iterative solution



# K-means clustering

- Iterative algorithm: Method 1
  1. Randomly initialize cluster means  $\mu_1, \mu_2, \dots, \mu_K$
  2. Compute distances  $D_{n1}, D_{n2}, \dots, D_{nK}$  for each object  $x_n$  to all  $K$  cluster means
  3. Assign the object  $x_n$  to cluster  $k$  with lowest distance  $D_{nk}$ 
    - i.e. assign  $x_n$  to its closest cluster  $k$  with mean  $\mu_k$
  4. Update each cluster mean  $\mu_k$ , to represent the mean of newly updated cluster


$$\mu_k = \frac{1}{N_k} \sum_{n=1}^{N_k} x_n$$

where  $N_k$  is the number of objects in  $k^{\text{th}}$  cluster

5. Stop if assignments don't change, or reached max number of iterations, else jump to step 2


# K-means clustering

- Randomly initialize cluster means  $\mu_1, \mu_2, \dots, \mu_k$
- Assign each  $x_n$  to its closest cluster with mean  $\mu_k$




# K-means clustering

- Update means  $\boldsymbol{\mu}_k = \frac{1}{N_k} \sum_{n=1}^{N_k} \mathbf{x}_n$




# K-means clustering

- Assign each  $x_n$  to its closest cluster with mean  $\mu_k$




# K-means clustering

- Update means  $\boldsymbol{\mu}_k = \frac{1}{N_k} \sum_{n=1}^{N_k} \mathbf{x}_n$




# K-means clustering

- Assign each  $x_n$  to its closest cluster with mean  $\mu_k$




# K-means clustering

- Update means  $\boldsymbol{\mu}_k = \frac{1}{N_k} \sum_{n=1}^{N_k} \mathbf{x}_n$



# K-means clustering


- Assign each  $x_n$  to its closest cluster with mean  $\mu_k$






# K-means clustering

- Update means  $\boldsymbol{\mu}_k = \frac{1}{N_k} \sum_{n=1}^{N_k} \mathbf{x}_n$




# K-means clustering

- Assign each  $x_n$  to its closest cluster with mean  $\mu_k$




# K-means clustering

- Update means  $\boldsymbol{\mu}_k = \frac{1}{N_k} \sum_{n=1}^{N_k} \mathbf{x}_n$




# K-means clustering

- Assign each  $x_n$  to its closest cluster with mean  $\mu_k$




# K-means clustering

- Update means  $\mu_k = \frac{1}{N_k} \sum_{n=1}^{N_k} \mathbf{x}_n$




# K-means clustering

- Assign each  $\mathbf{x}_n$  to its closest cluster with mean  $\boldsymbol{\mu}_k$




# K-means clustering

- Update means  $\boldsymbol{\mu}_k = \frac{1}{N_k} \sum_{n=1}^{N_k} \mathbf{x}_n$



# K-means clustering

- Solution at convergence





# K-means clustering

- Iterative algorithm: Method 2

1. Randomly assign each object  $\mathbf{x}_n$  to one of  $K$  clusters
2. Compute cluster means  $\boldsymbol{\mu}_1, \boldsymbol{\mu}_2, \dots, \boldsymbol{\mu}_K$  for each cluster  $k$ , to represent the mean of newly updated cluster


$$\boldsymbol{\mu}_k = \frac{1}{N_k} \sum_{n=1}^{N_k} \mathbf{x}_n$$

where  $N_k$  is the number of objects in  $k^{\text{th}}$  cluster

3. Compute distances  $D_{n1}, D_{n2}, \dots, D_{nK}$  for each object  $\mathbf{x}_n$  to all  $K$  cluster means
4. Assign the object  $\mathbf{x}_n$  to cluster  $k$  with lowest distance  $D_{nk}$ 
  - i.e. assign  $\mathbf{x}_n$  to its closest cluster  $k$  with mean  $\boldsymbol{\mu}_k$
5. Stop iterations if assignments don't change, or maximum number of iterations reached, else jump to step 2


# K-means clustering

- Cluster data points with Method 2




# K-means clustering

- Randomly assign each object  $x_n$  to one of  $K$  clusters




# K-means clustering

- Update means  $\mu_k = \frac{1}{N_k} \sum_{n=1}^{N_k} \mathbf{x}_n$




# K-means clustering

- Assign each  $x_n$  to its closest cluster with mean  $\mu_k$




# K-means clustering

- Update means  $\mu_k = \frac{1}{N_k} \sum_{n=1}^{N_k} \mathbf{x}_n$




# K-means clustering

- Assign each  $x_n$  to its closest cluster with mean  $\mu_k$



# K-means clustering

- Update means  $\mu_k = \frac{1}{N_k} \sum_{n=1}^{N_k} \mathbf{x}_n$





# K-means clustering


- The k-means clustering algorithm converges to a local minimum of the following *intra-group distance estimate*:

$$\mathcal{D} = \sum_{k=1}^K \sum_{n=1}^{N_k} (\mathbf{x}_n - \boldsymbol{\mu}_k)^T (\mathbf{x}_n - \boldsymbol{\mu}_k)$$

- i.e. total distance between all the objects and their respective cluster means
- Global optimum vs local optimum
  - Depends on random initialization
- Repeat runs of k-means clustering, pick one with lowest  $\mathcal{D}$



# How to choose $K$ ?

- Model selection
- 50 repeat runs
- $\mathcal{D}$  decreases with increasing  $K$ 
  - Why?
- $N = K$ ?
- $K$  often depends on application/problem/data




# K-means limitations

- K-means assumes clear cluster structure
  - What if it's lacking?
- Cluster means?
- Do objects conform to the similarity concept that K-means clustering heavily relies on?
- Outer cluster can't be represented by a single cluster mean



# Data transformation

- How about transforming data in to a new space where it can be separable?
  - $\mathbf{x} \rightarrow \phi(\mathbf{x})$
- For this example, consider  $\phi(\mathbf{x}) = x_n^{(1)2} + x_n^{(2)2}$ ?
- In practice, there is a very neat trick which doesn't even require the explicit data transformation
  - Kernel trick
- Kernel function (a kind of similarity measure)
  - A function that is equivalent to the dot product of vectors in the transformed space
  - $k(\mathbf{x}_m, \mathbf{x}_n) = \phi(\mathbf{x}_m)^T \phi(\mathbf{x}_n)$



# Kernel function

- There is a number of off-the-shelf kernels that have been shown to work well
- Linear kernel
  - $k(\mathbf{x}_m, \mathbf{x}_n) = \mathbf{x}_m^T \mathbf{x}_n$
- Gaussian kernel
  - $k(\mathbf{x}_m, \mathbf{x}_n) = \exp\{-\gamma(\mathbf{x}_m - \mathbf{x}_n)^T(\mathbf{x}_m - \mathbf{x}_n)\}$
  - $k(\mathbf{x}_m, \mathbf{x}_n) = \exp\{-\gamma\|\mathbf{x}_m - \mathbf{x}_n\|^2\}$
- Polynomial kernel
  - $k(\mathbf{x}_m, \mathbf{x}_n) = (\mathbf{x}_m^T \mathbf{x}_n + c)^\beta$

# Kernelized k-means

- Can we kernelize k-means?
- *Kernel function*
  - A function that corresponds to inner product of vectors in some other transformed space
- As long as an algorithm has data appearing only in inner products in model learning, kernels can be used

$$k(\mathbf{x}_m, \mathbf{x}_n) = \phi(\mathbf{x}_m)^T \phi(\mathbf{x}_n)$$

# Kernelized k-means

- Distance estimate (squared Euclidean)

$$D_{nk} = (\mathbf{x}_n - \boldsymbol{\mu}_k)^T (\mathbf{x}_n - \boldsymbol{\mu}_k)$$

- Cluster mean calculation

$$\boldsymbol{\mu}_k = \frac{1}{N_k} \sum_{n=1}^{N_k} \mathbf{x}_n$$

- Distance can be estimated as:

$$D_{nk} = \left( \mathbf{x}_n - \frac{1}{N_k} \sum_{m=1}^{N_k} \mathbf{x}_m \right)^T \left( \mathbf{x}_n - \frac{1}{N_k} \sum_{m=1}^{N_k} \mathbf{x}_m \right)$$

$$D_{nk} = \mathbf{x}_n^T \mathbf{x}_n - \frac{2}{N_k} \sum_{m=1}^{N_k} \mathbf{x}_m^T \mathbf{x}_n + \left( \frac{1}{N_k} \right)^2 \sum_{m=1}^{N_k} \sum_{l=1}^{N_k} \mathbf{x}_m^T \mathbf{x}_l$$

# Kernelized k-means

- Distance can be estimated as:

$$D_{nk} = \mathbf{x}_n^T \mathbf{x}_n - \frac{2}{N_k} \sum_{m=1}^{N_k} \mathbf{x}_m^T \mathbf{x}_n + \left( \frac{1}{N_k} \right)^2 \sum_{m=1}^{N_k} \sum_{l=1}^{N_k} \mathbf{x}_m^T \mathbf{x}_l$$

- With kernel trick, dot products could be replaced:

$$D_{nk} = k(\mathbf{x}_n, \mathbf{x}_n) - \frac{2}{N_k} \sum_{m=1}^{N_k} k(\mathbf{x}_m, \mathbf{x}_n) + \left( \frac{1}{N_k} \right)^2 \sum_{m=1}^{N_k} \sum_{l=1}^{N_k} k(\mathbf{x}_m, \mathbf{x}_l)$$

- Let's recall  $\boldsymbol{\mu}_k$  will be

$$\boldsymbol{\mu}_k = \frac{1}{N_k} \sum_{n=1}^{N_k} \phi(\mathbf{x}_n)$$

but what's  $\phi(\mathbf{x}_n)$ ?



# Kernelized k-means


- Algorithm:

1. Randomly assign each object  $\mathbf{x}_n$  to one of  $K$  clusters
2. Compute distances  $D_{n1}, D_{n2}, \dots, D_{nK}$  for each object  $\mathbf{x}_n$ , using the kernel trick


$$D_{nk} = k(\mathbf{x}_n, \mathbf{x}_n) - \frac{2}{N_k} \sum_{m=1}^{N_k} k(\mathbf{x}_m, \mathbf{x}_n) + \left(\frac{1}{N_k}\right)^2 \sum_{m=1}^{N_k} \sum_{l=1}^{N_k} k(\mathbf{x}_m, \mathbf{x}_l)$$

3. Assign the object  $\mathbf{x}_n$  to cluster  $k$  with lowest distance  $D_{nk}$ 
  - i.e. assign  $\mathbf{x}_n$  to its closest cluster
4. Stop iterations if assignments don't change, or maximum number of iterations reached, else jump to step 2


# Kernelized k-means




(a) Kernel  $K$ -means after one iteration



(b) After five iterations



(c) After 10 iterations



(d) At convergence (30 iterations)

# Kernelized k-means

- Brings flexibility to simple k-means algorithm
- Need to set additional kernel parameters
- Sensitive to initialization
  - Often, it's good to run standard k-means algorithm before kernelized k-means for decent initialization

# K-means: limitations

- Sensitive to initialization
- Computationally expensive
- Often needs repeated runs
- Makes hard assignments (i.e. each object can belong only to one cluster)
  - Soft assignment (each object has a probability of belonging to each cluster)

# Hierarchical clustering

- Bottom-up (agglomerative) approach
  - Each object starts as a singleton cluster
  - Two most similar clusters are merged iteratively
  - Stop when all objects are in same cluster
- Top-down (divisive) approach
  - All objects belong to same cluster
  - “outsider” objects from least cohesive cluster are removed iteratively
  - Stop when each object is a singleton cluster

# Clustering: similarity measure

- We know how to estimate similarity (or distance) between pair of objects
- How do we measure similarity between a pair of clusters ( $R$  and  $S$ )?




- Linkage
  - Min/single link
  - Max/complete link
  - Average link

# Clustering: similarity measure

- Min/single link:

- Clusters can get very large


$$D(R, S) = \min_{p \in R, q \in S} D(p, q)$$



- Max/complete link:

- Results in small, round clusters


$$D(R, S) = \max_{p \in R, q \in S} D(p, q)$$



- Average link:

- Compromise between min and max links

$$D(R, S) = \frac{1}{|R||S|} \sum_{p \in R, q \in S} D(p, q)$$




# Hierarchical clustering

- Hierarchical agglomerative clustering (HAC) is a bottom-up clustering strategy
    - Each object is a singleton cluster
    - Successively merge closest pair of clusters together, till all clusters merge
1. Consider each object as a singleton cluster
  2. Compute distance between all pairs of clusters
  3. Remove the closest pair of clusters from the data, and replace them as a single cluster
  4. If not all objects belong to same single cluster, jump to step 2





# Hierarchical clustering

- Hierarchically cluster



# Hierarchical clustering

- The results of hierarchical clustering can be visualized as a dendrogram where the height of each “junction” in the dendrogram represents the distance between the pair of clusters
- Hierarchical clustering generates all number of clusters
  - Cut horizontally across the dendrogram



# Flat clustering vs hierarchical clustering

- Flat clustering (K-means) results in a single partitioning of objects
- Hierarchical clustering can produce different partitioning results depending on the level-of-resolution we are looking at
- Flat clustering needs to know  $K$  in advance
- Hierarchical clustering doesn't need to know  $K$

# Flat clustering vs hierarchical clustering

- Flat clustering is usually more efficient (run-time) – if we know  $K$ 
  - $N \times K$  distance computations at each iteration
- Hierarchical clustering can be slow (has to make several merge/split decisions)
  - Choosing 2 closest objects from  $N$  objects requires  $\frac{N!}{2!(N-2)!} \Rightarrow N * (N - 1)/2$  calculations just for first pairing
- No clear consensus on which of the two produces better clustering
  - Results depend upon nature of data

# Summary

- Two common clustering methods
  - By associating an object with its nearest cluster mean
  - By successively merging smaller clusters to hierarchically form larger clusters
- Free choices
  - How many clusters?
  - What distance metric?
  - What linkage method?
- Optimal combination is often found by trial-and-error
- Kernelized k-means

# Exercise (ungraded)

- Use K-means algorithm and Euclidean distance metric to cluster the following 2-dimensional objects in to 3 clusters
  - O1 (2,10), O2 (2,5), O3 (8,4), O4 (5,8), O5 (7,5), O6 (6,4), O7 (1,2), O8 (4,9)
- Suppose that initial cluster centres are O1, O4, and O7. Run the K-means algorithm steps for 1 iteration and show:
  - The clusters (i.e. objects belonging to each cluster)
  - The centres of new clusters
  - Draw a 10x10 grid with all the 8 objects and show the clusters and centres after the first iteration

# Exercise (ungraded)

- Use min/single link to perform agglomerative clustering by showing the dendrogram

- The data is described by the distance matrix

	A	B	C	D
A	0	1	4	5
B		0	2	6
C			0	3
D				0

- Use max/complete link to perform agglomerative clustering by showing the dendrogram
  - The data is described by the distance matrix
- Note: the height of each “junction” in the dendrogram represents the distance between the pair of clusters


# Exercise (ungraded)

- Use single and complete link agglomerative clustering to cluster the following 8 objects by showing the dendrograms
  - O1 (2,10), O2 (2,5), O3 (8,4), O4 (5,8), O5 (7,5), O6 (6,4), O7 (1,2), O8 (4,9)



# Exercise (ungraded)

- Try MATLAB code - kmeansexample.m (from FCML book website)
- Try MATLAB code - kmeansK.m (from FCML book website)
- Try MATLAB code - kernelkmeans.m (from FCML book website)
- Try MATLAB code – kmeans\_cluster.m (from Canvas)
- Try MATLAB code – run\_kmeans.m (from Canvas)



Author's material  
(Simon Rogers)



Iain Styles



Thank You