

# Machine Learning, Machine Learning (extended)

## 10 – Supervised Learning: Ensemble Methods

Kashif Rajpoot

[k.m.rajpoot@cs.bham.ac.uk](mailto:k.m.rajpoot@cs.bham.ac.uk)

School of Computer Science  
University of Birmingham

# Outline

- Ensemble methods
  - Boosting
  - Bagging
- Decision tree
- Random forests

# Ensemble methods

- Combining ‘weak classifiers’ in order to produce a ‘strong classifier’
  - “two heads are better than one”
- **Boosting**: train a new classifier focusing on training samples misclassified by an earlier classifier
  - Weak classifier: any classifier better than a random guess
  - AdaBoost
- **Bagging** (**b**ootstrap **agg**regation): generate new training data as a random subset of original data and train a new classifier on this subset
  - Weak classifier: a decision tree classifier
  - Random forests

# Modified from Randomized Forests for Visual Recognition

Jamie Shotton

Tae-Kyun Kim

Björn Stenger



**TOSHIBA**

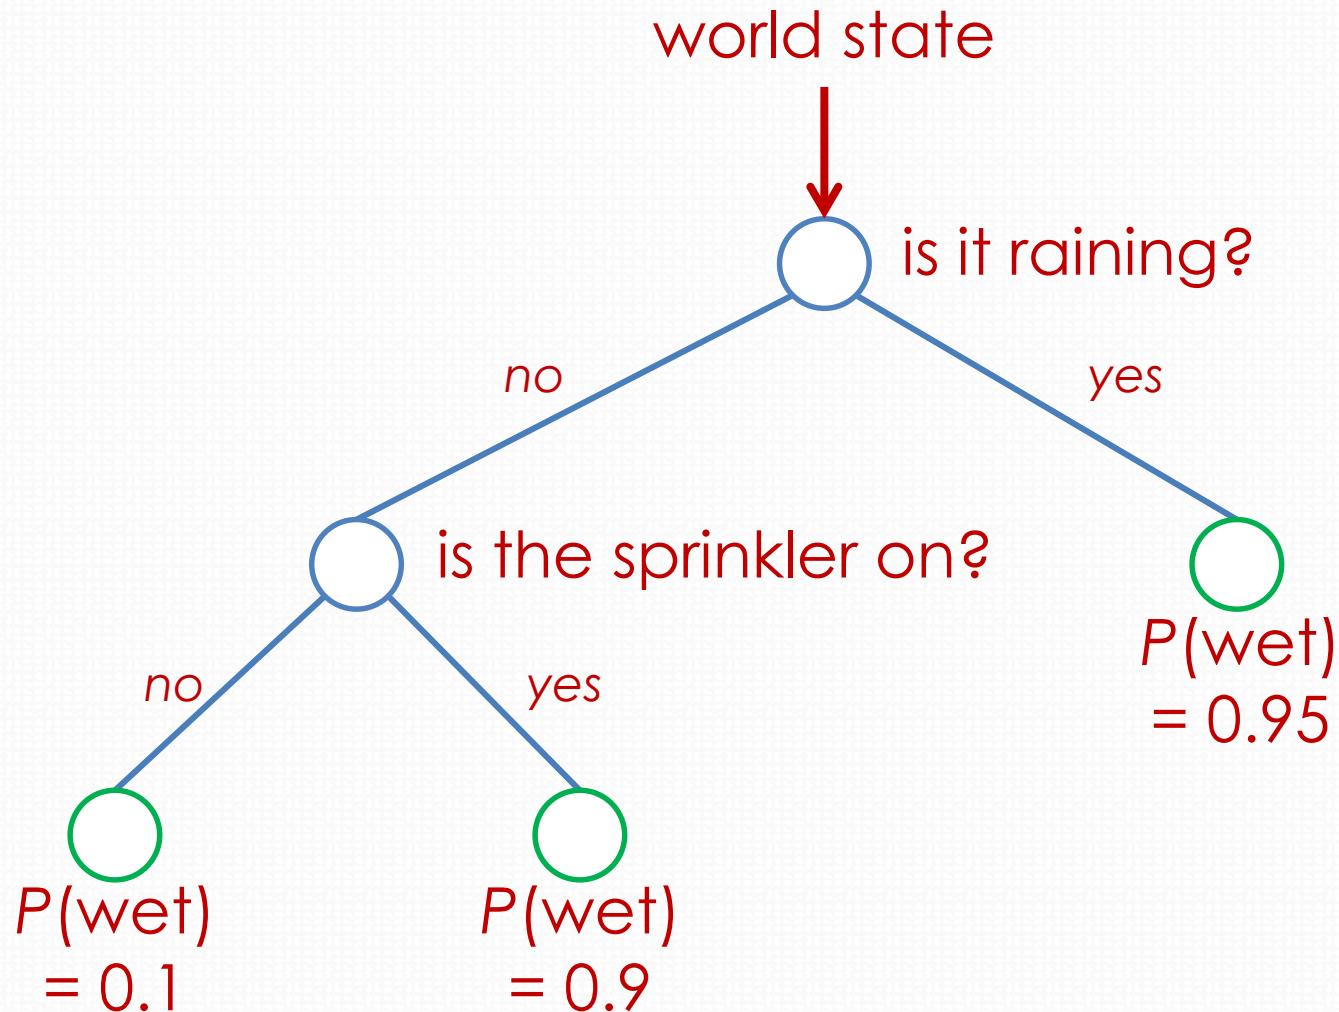
ICCV 2009, Kyoto, Japan

# Randomized decision forests

- Very fast tool for classification
- Good generalization through randomized training
- Inherently multi-class
- Simple training / testing algorithms

“Randomized Decision Forests” = “Randomized Forests” =  
“Random Forests<sup>TM</sup>”

# Basics: is the grass wet?



# Basics: binary decision tree

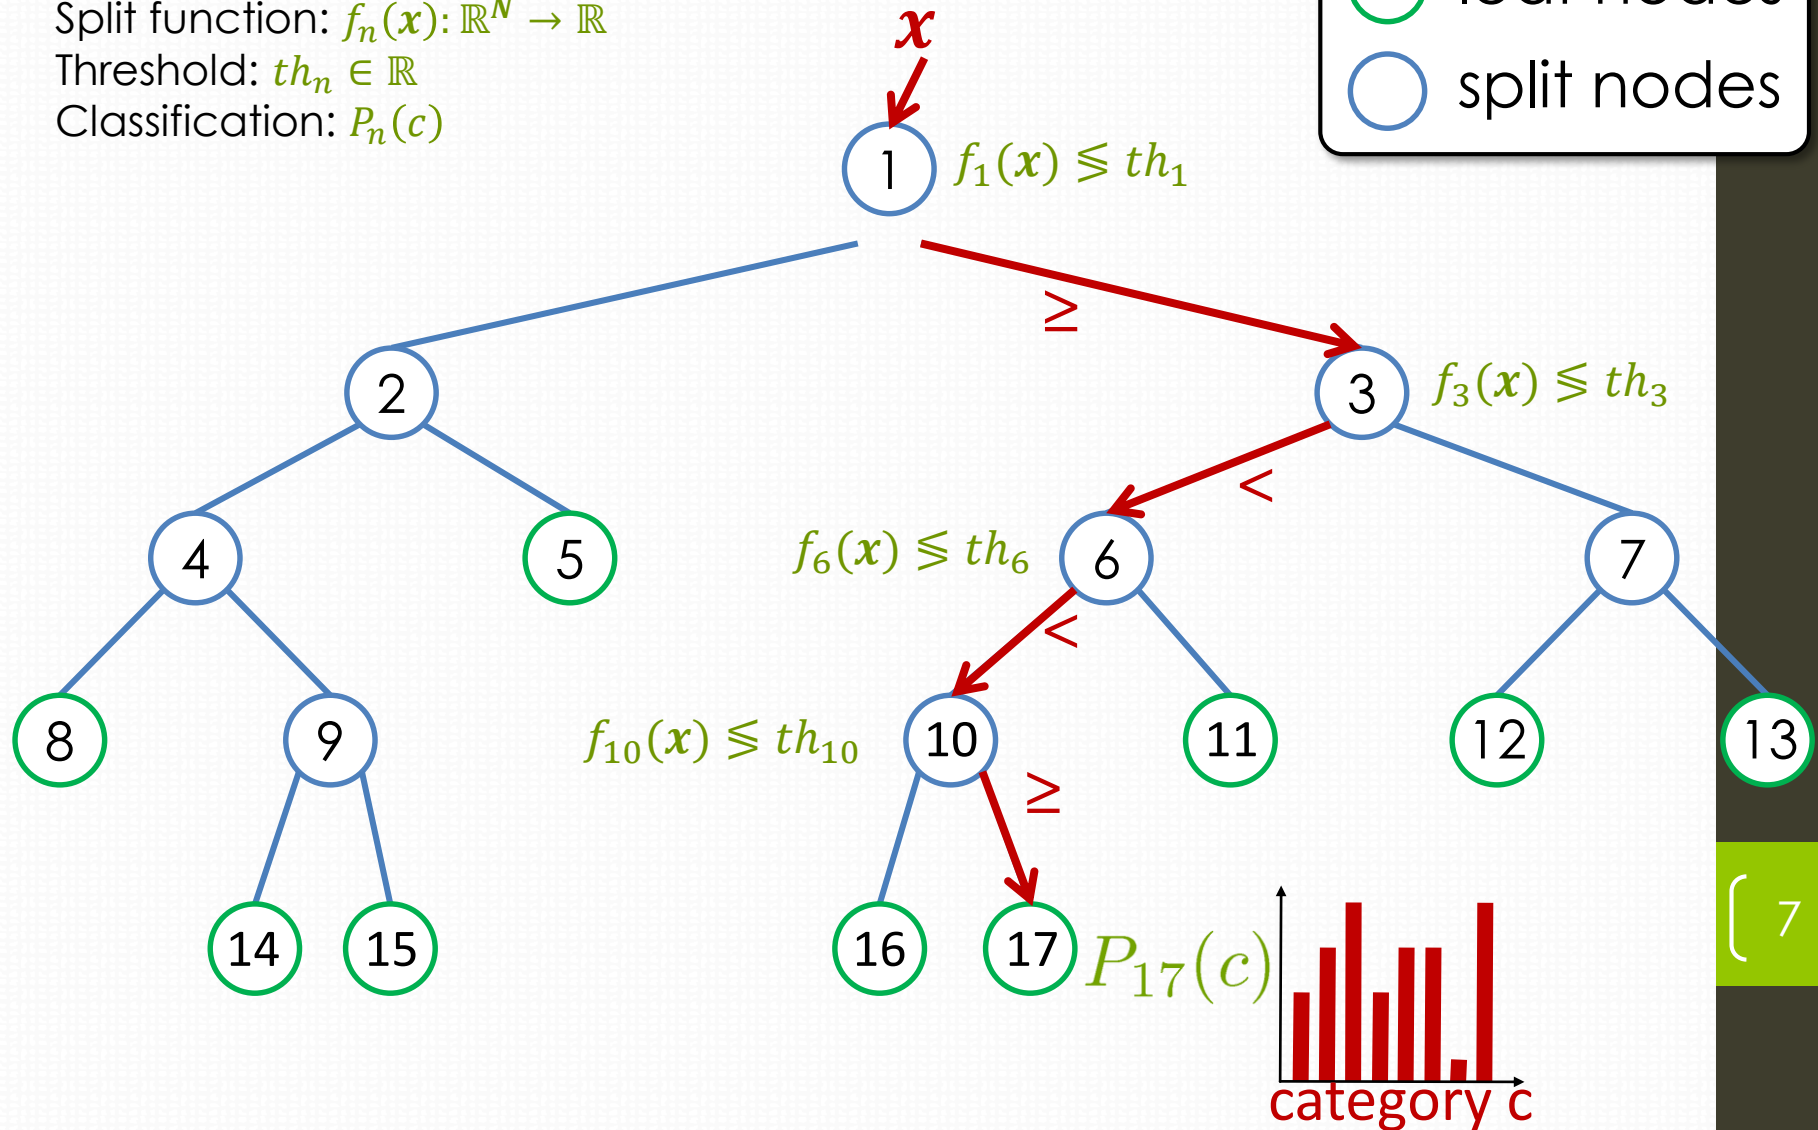
Feature/attribute vector:  $\mathbf{x} \in \mathbb{R}^N$

Split function:  $f_n(\mathbf{x}): \mathbb{R}^N \rightarrow \mathbb{R}$

Threshold:  $th_n \in \mathbb{R}$

Classification:  $P_n(c)$

○ leaf nodes  
○ split nodes




# Decision tree classification: pseudo-code

```
double[] ClassifyDT(node, x)
  if node.IsSplitNode then
    if node.f(x) >= node.th then
      return ClassifyDT(node.right, x)
    else
      return ClassifyDT(node.left, x)
    end
  else
    return node.P
  end
end
```



# Learning example

- Try several lines, 'chosen at random'
- Keep line that best separates data
  - Maximize information gain
- Recurse




- feature vectors are  $x, y$  coordinates:
- split functions are lines with parameters  $a, b$ :
- threshold determines intercepts:
- four classes: purple, blue, red, green

$$\begin{aligned} \mathbf{x} &= [x, y]^T \\ f_n(\mathbf{x}) &= ax + by \\ th_n \end{aligned}$$

# Learning example

- Try several lines, 'chosen at random'
- Keep line that best separates data
  - Maximize information gain
- Recurse




- feature vectors are  $x, y$  coordinates:
- split functions are lines with parameters  $a, b$ :
- threshold determines intercepts:
- four classes: purple, blue, red, green

$$\begin{aligned} \mathbf{x} &= [x, y]^T \\ f_n(\mathbf{x}) &= ax + by \\ th_n \end{aligned}$$

# Learning example

- Try several lines, 'chosen at random'
- Keep line that best separates data
  - Maximize information gain
- Recurse




- feature vectors are  $x, y$  coordinates:
- split functions are lines with parameters  $a, b$ :
- threshold determines intercepts:
- four classes: purple, blue, red, green

$$\begin{aligned}\mathbf{x} &= [x, y]^T \\ f_n(\mathbf{x}) &= ax + by \\ th_n\end{aligned}$$

# Learning example

- Try several lines, 'chosen at random'
- Keep line that best separates data
  - Maximize information gain
- Recurse



- feature vectors are  $x, y$  coordinates:
- split functions are lines with parameters  $a, b$ :
- threshold determines intercepts:
- four classes: purple, blue, red, green

$$\begin{aligned} \mathbf{x} &= [x, y]^T \\ f_n(\mathbf{x}) &= ax + by \\ th_n \end{aligned}$$

# Randomized learning

- Randomness in:
  - Bagging: randomly select the subset of data at the root of a tree
  - Randomly select the features at a tree node
    - “feature = attribute” in machine learning
  - Randomly select the threshold value

# Randomized learning

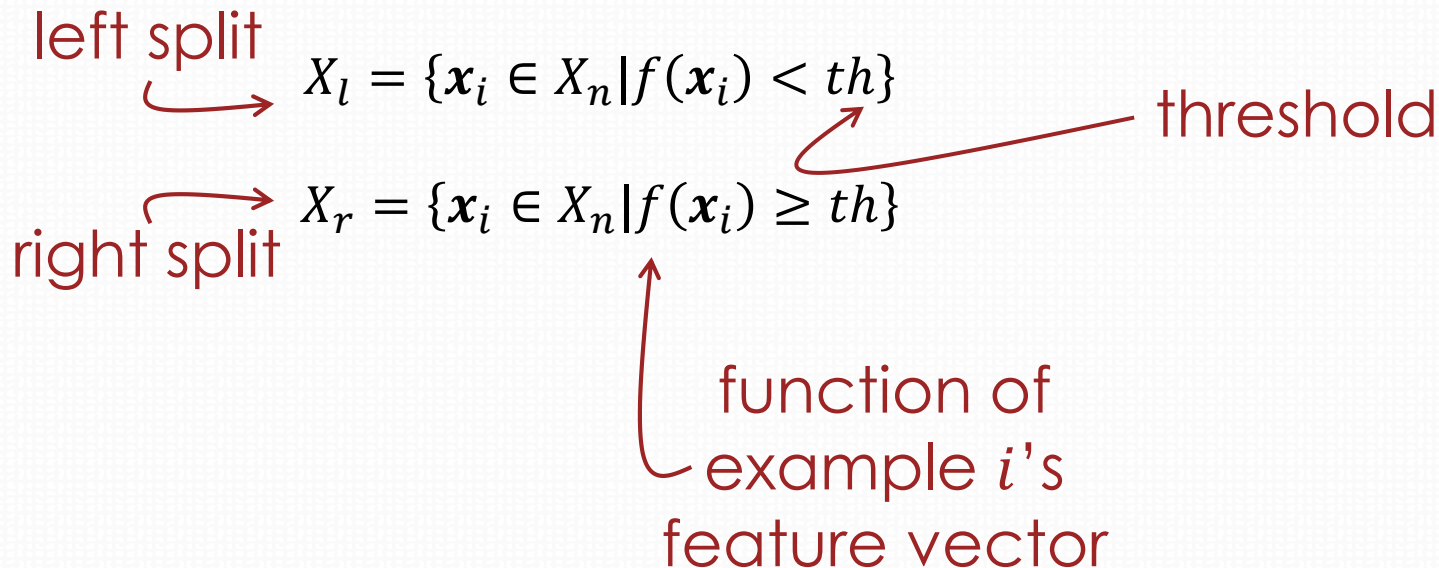
- Randomly select  $X_n$  examples as a subset from  $X$  examples
- Recursively split  $X_n$  examples at node  $n$

left split  $\rightarrow X_l = \{\mathbf{x}_i \in X_n | f(\mathbf{x}_i) < th\}$

right split  $\rightarrow X_r = \{\mathbf{x}_i \in X_n | f(\mathbf{x}_i) \geq th\}$

threshold

function of example  $i$ 's feature vector



# Randomized learning

- Features  $f(\mathbf{x})$  chosen at random from feature pool  $F$
- Threshold  $th$  chosen at random in range
  - $th \in (\min(f(\mathbf{x})), \max(f(\mathbf{x})))$
- Choose  $f$  and  $th$  to maximize an objective function (e.g. information gain, Gini index)
  - Estimates whether it's “good” to distribute data further

# Randomized learning

- $P_n(c)$  is the histogram (i.e. count) of example labels of class  $c$  which reached node  $n$
- For example, at a leaf node, if 200 training example reach and there are 3 classes with following count, then the  $P_n(c)$  is estimated as:

	$c = 1$	$c = 2$	$c = 3$
Count of examples	12	134	54
$P_n(c)$	12/200	134/200	54/200
$P_n(c)$	0.06	0.67	0.27



# Implementation details

- How many features and thresholds to try?
  - just one = “extremely randomized”
  - few  $\rightarrow$  fast training, may under-fit
  - many  $\rightarrow$  slower training, may over-fit
- When to stop growing the tree?
  - maximum depth
  - minimum information gain

# Decision tree learning: pseudo-code

```
TreeNode LearnDT(X)
  repeat featureTests times
    let f = RndFeature()
    let r = EvaluateFeatureResponses(X, f)

    repeat threshTests times
      let th = RndThreshold(r)
      let (X_l, X_r) = Split(X, r, th)

      let gain = InfoGain(X_l, X_r)
      if gain is best then remember f, th, X_l, X_r
    end
  end

  if best gain is sufficient
    return SplitNode(f, th, LearnDT(X_l), LearnDT(X_r))
  else
    return LeafNode(HistogramExamples(X_s))
  end
end
```

# Binary decision tree: summary

- Fast greedy training algorithm
  - can search infinite pool of features
  - heterogeneous pool of features
- Fast testing algorithm
- Needs careful choice of hyper-parameters
  - maximum depth
  - number of features and thresholds to try

# Information gain and entropy

- Information gain: gain in information (i.e purity of data according to class labels) by split of data from parent to child nodes in the tree

- $$IG(f_n) = E(parent) - \frac{|X_l|}{|X_n|} * E(left) - \frac{|X_r|}{|X_n|} * E(right)$$


where  $|X_n|$  denotes number of data samples at node  $n$

- Entropy: measure of disorder (or impurity) in a bunch of data samples

- $$E = - \sum_{c=1}^C P_c \log_2(P_c)$$



Low entropy



High entropy

# Information gain and entropy

- $E = -\sum_{c=1}^C P_c \log_2(P_c) = ?$ 
  - $E = -P_{Flu=Y} \log_2(P_{Flu=Y}) - P_{Flu=N} \log_2(P_{Flu=N}) = ?$
  - $E = -\frac{5}{8} \log_2\left(\frac{5}{8}\right) - \frac{3}{8} \log_2\left(\frac{3}{8}\right) = 0.9544$

chills	runny nose	headache	fever	Flu?
Y	N	Mild	Y	N
Y	Y	No	N	Y
Y	N	Strong	Y	Y
N	Y	Mild	Y	Y
N	N	No	N	N
N	Y	Strong	Y	Y
N	Y	Strong	N	N
Y	Y	Mild	Y	Y

# Information gain and entropy

$$E(\text{left}) = ?$$

$$E(\text{chills} = Y) = -P_{Flu=Y} \log_2(P_{Flu=Y}) - P_{Flu=N} \log_2(P_{Flu=N}) = ?$$

$$E(\text{chills} = Y) = -\frac{3}{4} \log_2\left(\frac{3}{4}\right) - \frac{1}{4} \log_2\left(\frac{1}{4}\right) = 0.8113$$

$$IG(\text{chills}) = E(\text{parent}) - \frac{|X_l|}{|X_n|} * E(\text{left}) - \frac{|X_r|}{|X_n|} * E(\text{right}) = ?$$

$$IG(\text{chills}) = 0.9544 - 0.5 * 0.8113 - 0.5 * 1 = 0.0488$$

$$E(\text{right}) = ?$$

$$E(\text{chills} = N) = -P_{Flu=Y} \log_2(P_{Flu=Y}) - P_{Flu=N} \log_2(P_{Flu=N}) = ?$$

$$E(\text{chills} = N) = -\frac{2}{4} \log_2\left(\frac{2}{4}\right) - \frac{2}{4} \log_2\left(\frac{2}{4}\right) = 1$$

$$\frac{|X_l|}{|X_n|} = ?$$

$$\frac{|X_{\text{chills}=Y}|}{|X_n|} = ?$$

$$\frac{|X_{\text{chills}=Y}|}{|X_n|} = \frac{4}{8}$$

$$\frac{|X_r|}{|X_n|} = ?$$

$$\frac{|X_{\text{chills}=N}|}{|X_n|} = ?$$

$$\frac{|X_{\text{chills}=N}|}{|X_n|} = \frac{4}{8}$$

chills	runny nose	headache	fever	Flu?
Y	N	Mild	Y	N
Y	Y	No	N	Y
Y	N	Strong	Y	Y
N	Y	Mild	Y	Y
N	N	No	N	N
N	Y	Strong	Y	Y
N	Y	Strong	N	N
Y	Y	Mild	Y	Y

# Information gain and entropy

$$E(\text{left}) = ?$$

$$E(\text{runny} = Y) = -P_{Flu=Y} \log_2(P_{Flu=Y}) - P_{Flu=N} \log_2(P_{Flu=N}) = ?$$

$$E(\text{runny} = Y) = -\frac{4}{5} \log_2\left(\frac{4}{5}\right) - \frac{1}{5} \log_2\left(\frac{1}{5}\right) = 0.7219$$

$$IG(\text{runny}) = E(\text{parent}) - \frac{|X_l|}{|X_n|} * E(\text{left}) - \frac{|X_r|}{|X_n|} * E(\text{right}) = ?$$

$$IG(\text{runny}) = 0.9544 - 0.625 * 0.7219 - 0.375 * 0.9183 = 0.1589$$

$$E(\text{right}) = ?$$

$$E(\text{runny} = N) = -P_{Flu=Y} \log_2(P_{Flu=Y}) - P_{Flu=N} \log_2(P_{Flu=N}) = ?$$

$$E(\text{runny} = N) = -\frac{1}{3} \log_2\left(\frac{1}{3}\right) - \frac{2}{3} \log_2\left(\frac{2}{3}\right) = 0.9183$$

$$\frac{|X_l|}{|X_n|} = ?$$

$$\frac{|X_{\text{runny}=Y}|}{|X_n|} = ?$$

$$\frac{|X_{\text{runny}=Y}|}{|X_n|} = \frac{5}{8}$$

$$\frac{|X_r|}{|X_n|} = ?$$

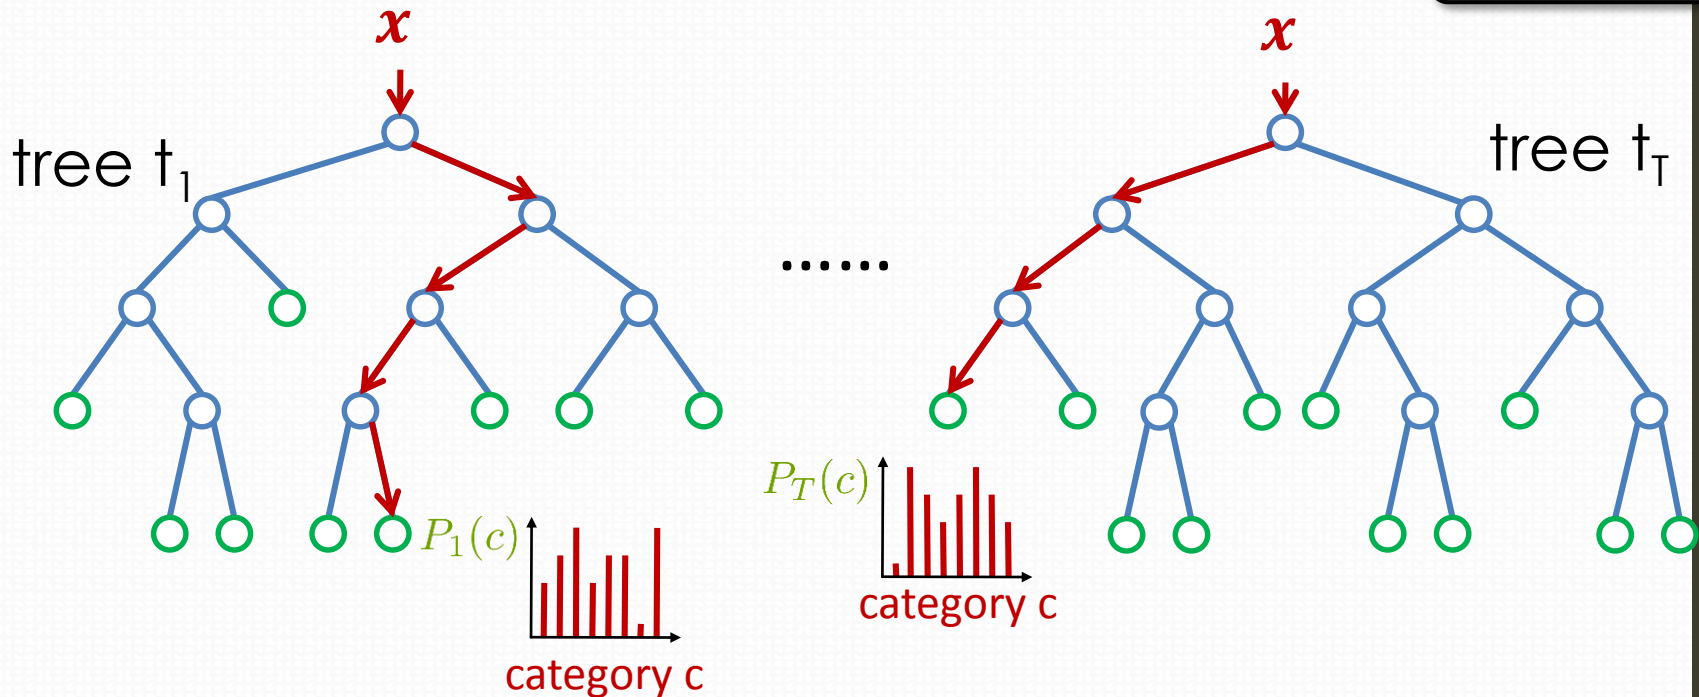
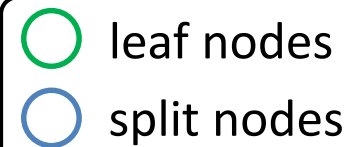
$$\frac{|X_{\text{runny}=N}|}{|X_n|} = ?$$

$$\frac{|X_{\text{runny}=N}|}{|X_n|} = \frac{3}{8}$$

chills	runny nose	headache	fever	Flu?
Y	N	Mild	Y	N
Y	Y	No	N	Y
Y	N	Strong	Y	Y
N	Y	Mild	Y	Y
N	N	No	N	N
N	Y	Strong	Y	Y
N	Y	Strong	N	N
Y	Y	Mild	Y	Y

# A forest of trees

- Forest is ensemble of several decision trees



- Classification

- $$P(c|\mathbf{x}) = \frac{1}{T} \sum_{tr=1}^T P_{tr}(c|\mathbf{x})$$



# Decision forests: pseudo-code

```
double[] ClassifyDF(forest, x)
    // allocate memory
    let P = double[forest.CountClasses]

    // loop over trees in forest
    for tr = 1 to forest.CountTrees
        let P' = ClassifyDT(forest.Tree[tr], x)
        P = P + P' // sum distributions
    end

    // normalise
    P = P / forest.CountTrees
end
```

# Learning a forest

- Divide training examples into  $T$  subsets
  - $X_{tr} \subseteq X$
  - improves generalization
  - reduces memory requirements & training time
- Subsets are chosen at random
- Subsets can have overlap (and usually do)
- Train each decision tree  $tr$  on subset  $X_{tr}$

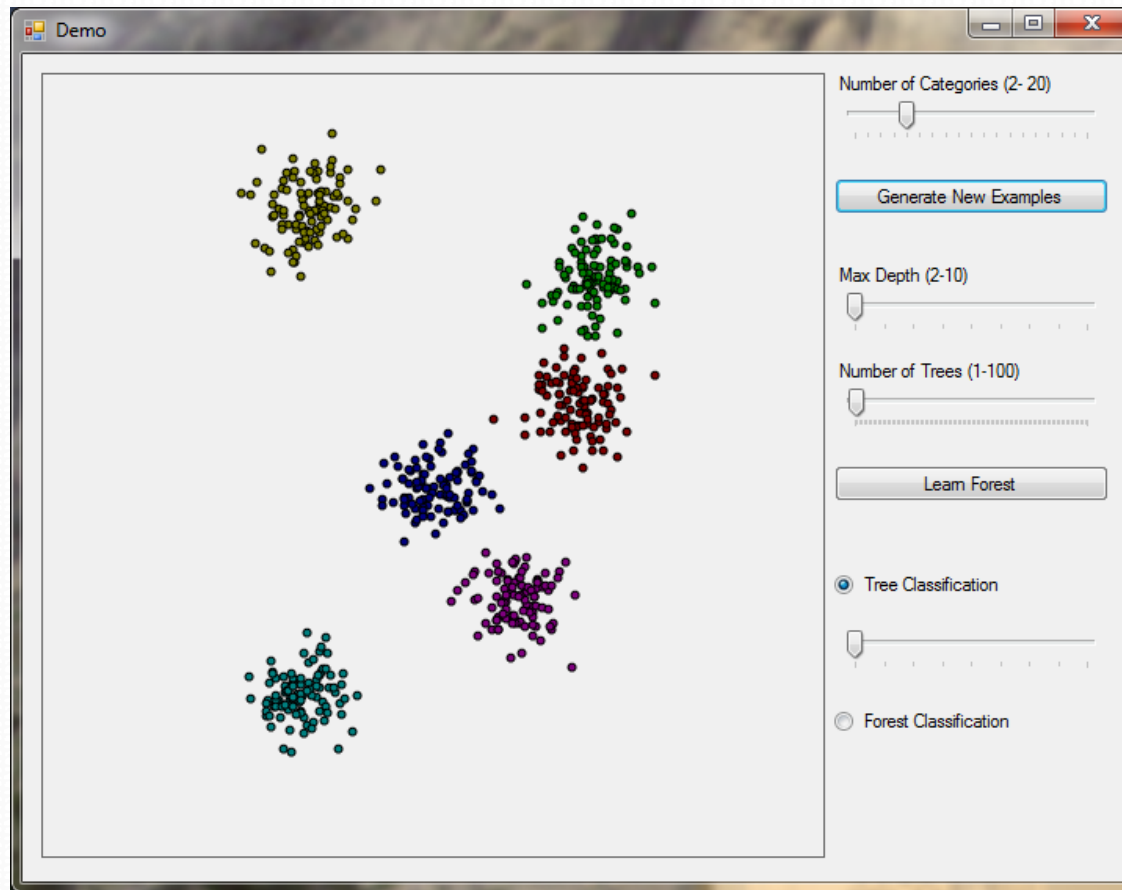
# Learning a forest: pseudo-code

```
Forest LearnDF(countTrees, X)
  // allocate memory
  let forest = Forest(countTrees)

  // loop over trees in forest
  for tr = 1 to countTrees
    let X_tr = RandomSplit(X)
    forest[tr] = LearnDT(X_tr)
  end

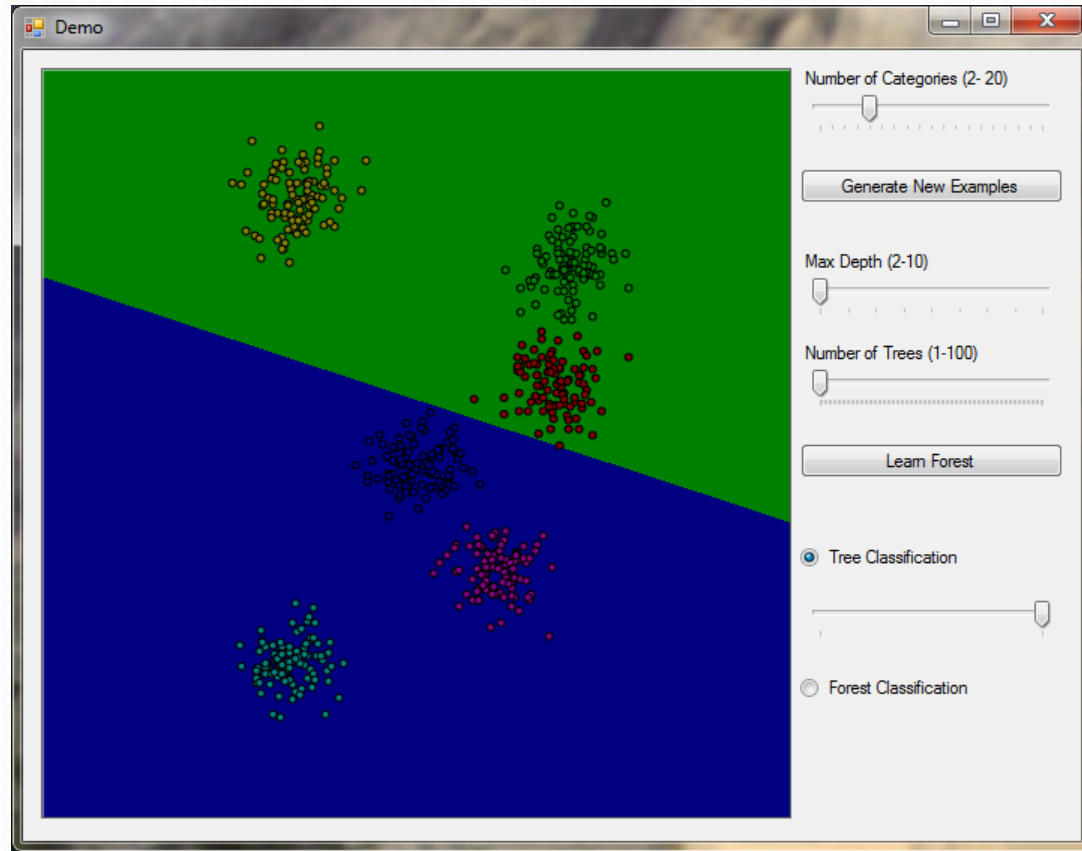
  // return forest object
  return forest
end
```

# Random forest: demo



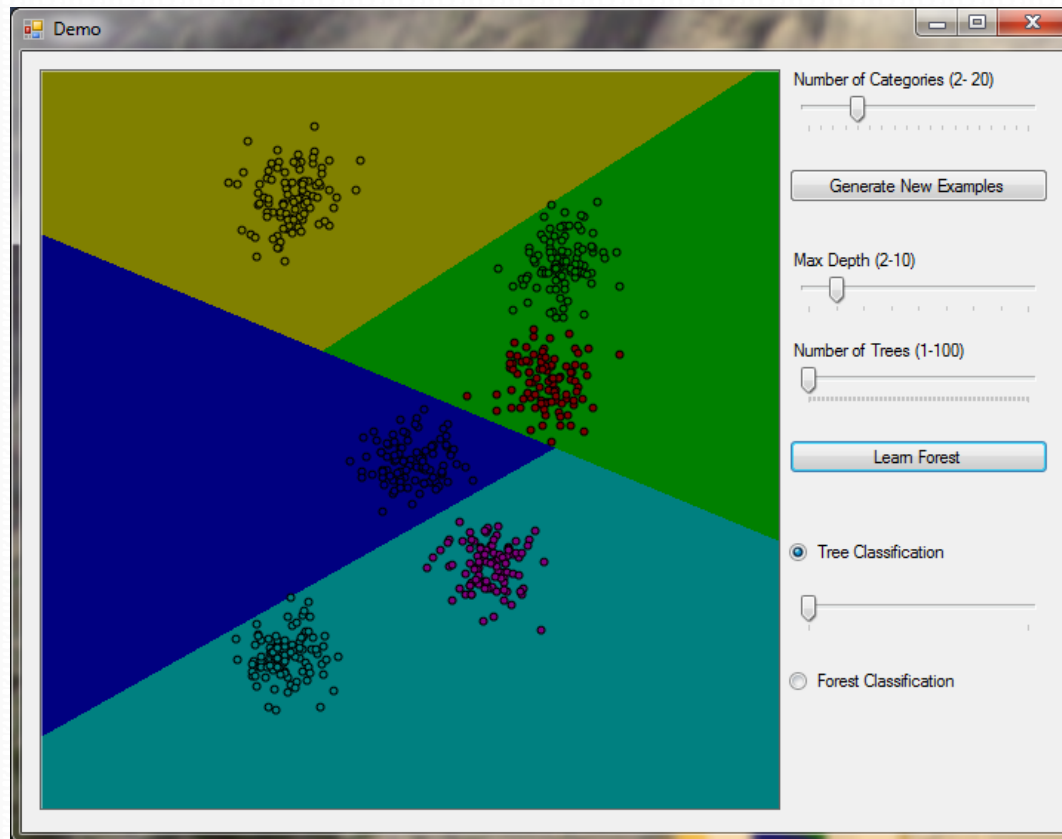
6 classes in a 2 dimensional feature space.  
Split functions are lines in this space.

# Random forest: demo



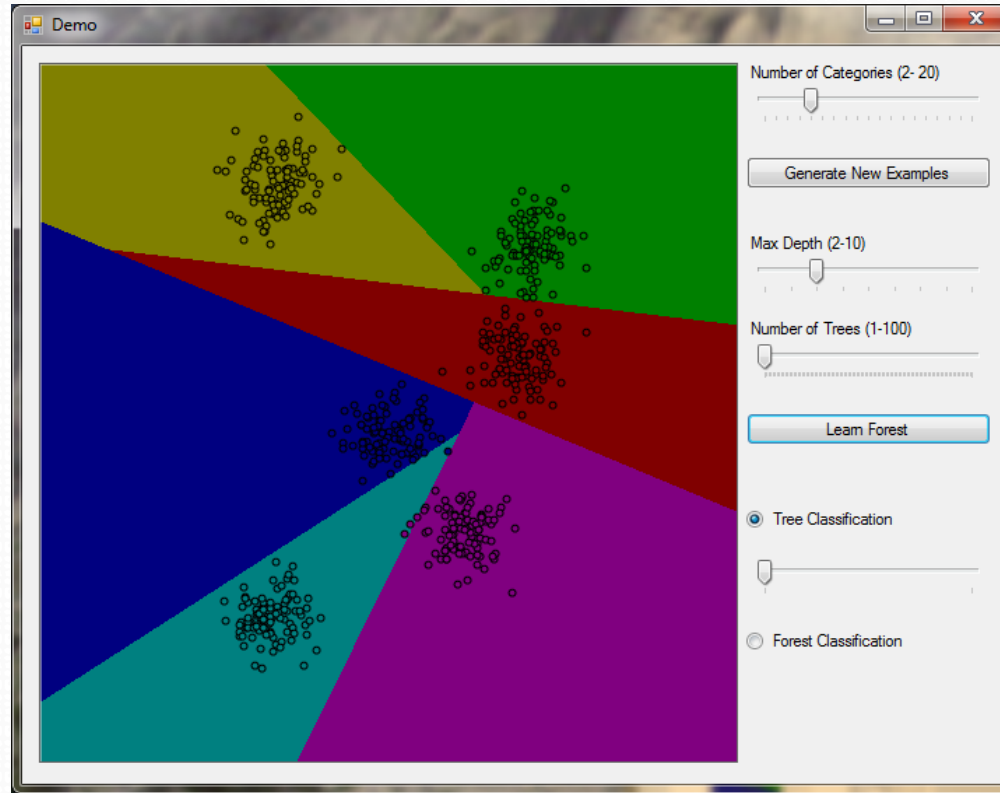
With a depth 2 tree, you cannot separate all six classes.

# Random forest: demo



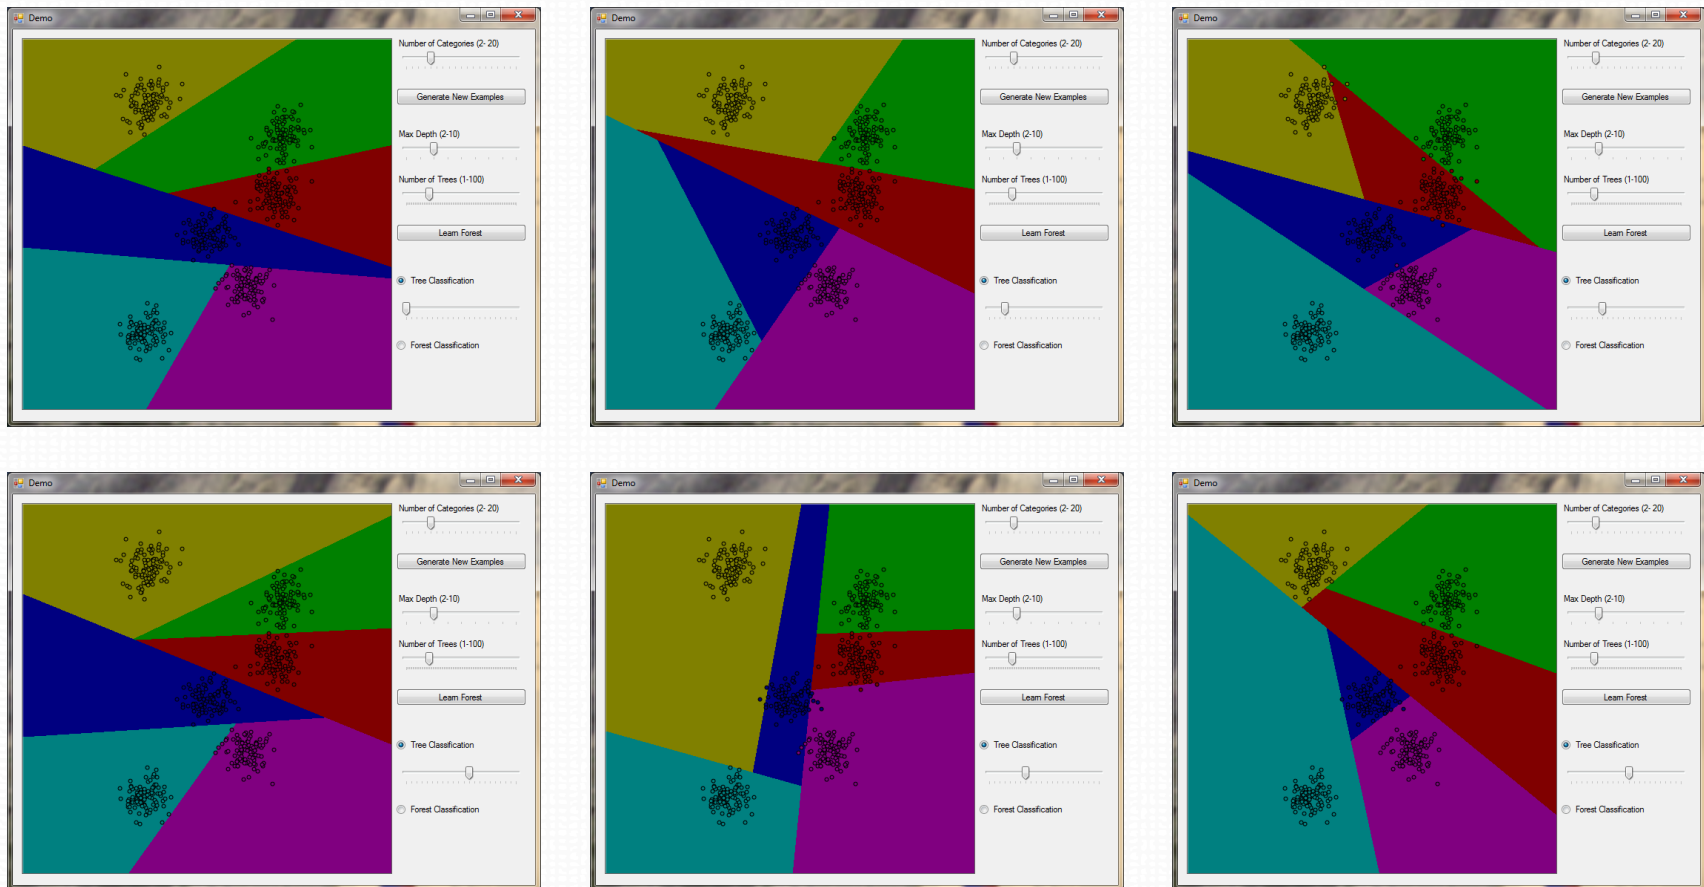
With a depth 3 tree, you are doing better, but still cannot separate all six classes.

# Random forest: demo



With a depth 4 tree, you now have at least as many leaf nodes as classes, and so are able to classify most examples correctly.

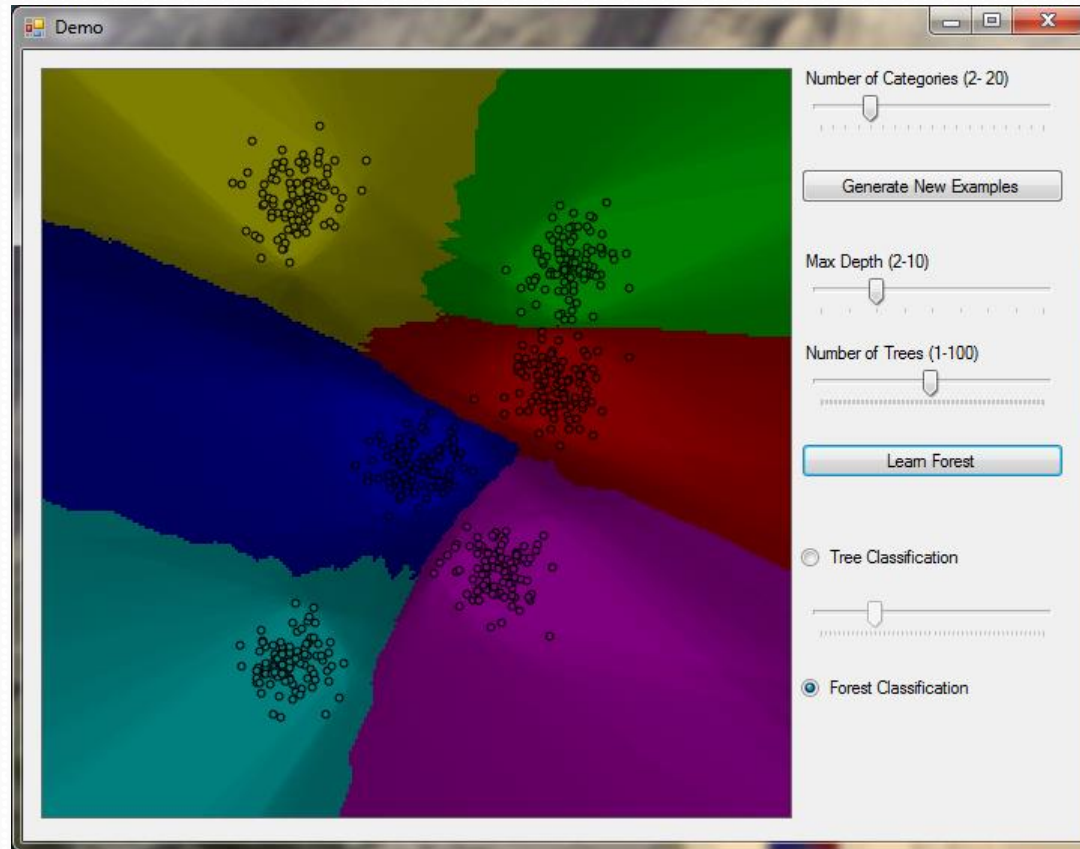
# Random forest: demo



Different trees within a forest can give rise to very different decision boundaries, none of which is particularly good on its own.



# Random forest: demo



But averaging together many trees in a forest can result in decision boundaries that look very sensible, and are even quite close to the max margin classifier.

# Summary

- Very fast classification algorithm
- Accuracy comparable with other classifiers
- Simple to implement

# Further reading/References

- ICCV'2009 Tutorial
  - <http://jamie.shotton.org/work/presentations/ICCV2009TutorialPartI.pptx>
- Random Forests for Regression and Classification; by Adele Cutler
  - <http://www.math.usu.edu/adele/RandomForests/Ovronnaz.pdf>
- Machine Learning; by Tom Mitchell (Chapter 3)
- Pattern Classification; By Duda, Hart, Stork (Chapter 8)



Thank You